

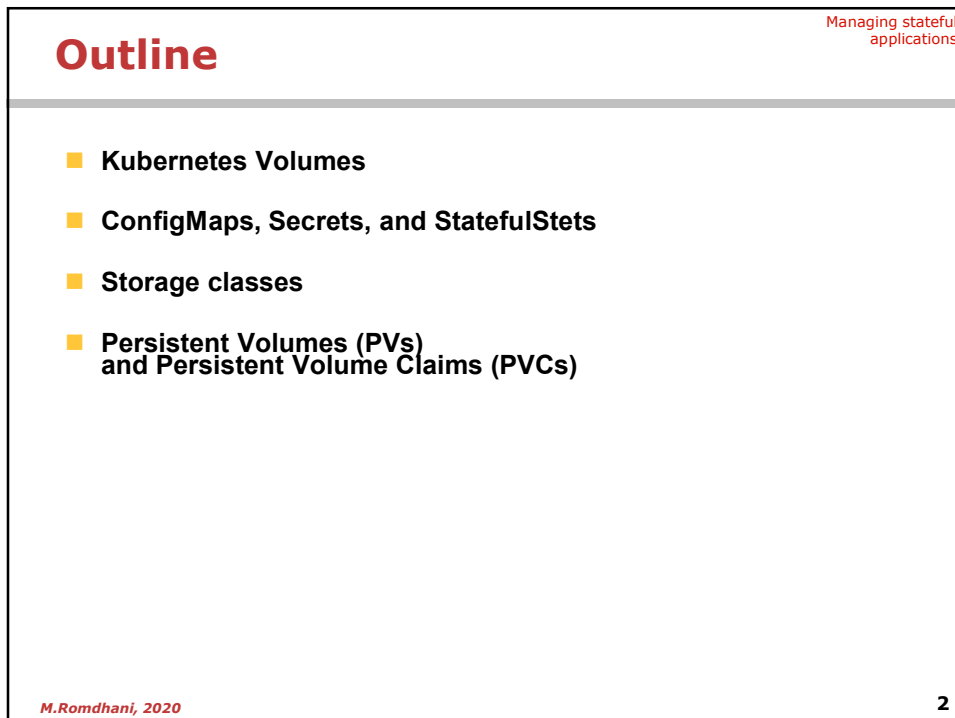
The slide features a purple header and footer with a collage of images. In the center, the Kubernetes logo is displayed above the text "Unit 4". Below this, the main title "Managing Stateful applications in Kubernetes" is written in a large, bold, red font. In the bottom right corner, there are three small icons (a circle, a square, and a triangle) above the text "Business Training".

Unit 4

Managing Stateful applications in Kubernetes

Business Training

1



The slide has a purple header with the text "Managing stateful applications" in red. Below the header, the word "Outline" is written in a large, bold, red font. A list of four topics is presented, each preceded by a yellow square bullet point. At the bottom left, the text "M.Romdhani, 2020" is visible, and at the bottom right, the number "2" is displayed.

Outline

Managing stateful applications

- Kubernetes Volumes
- ConfigMaps, Secrets, and StatefulSets
- Storage classes
- Persistent Volumes (PVs) and Persistent Volume Claims (PVCs)

M.Romdhani, 2020

2

2

Kubernetes Volumes

3

Volumes

Managing stateful
applications

- **Volumes are special directories that are mounted in containers**
- **Volumes can have many different purposes:**
 - share files and directories between containers running on the same machine
 - share files and directories between containers and their host
 - centralize configuration information in Kubernetes and expose it to containers
 - manage credentials and secrets and expose them securely to containers
 - store persistent data for stateful services
 - access storage systems (like EBS, NFS, Portworx, and many others)
- **Kubernetes volumes vs. Docker volumes**
 - Kubernetes and Docker volumes are very similar
 - Docker volumes allow us to share data between containers running on the same host
 - Kubernetes volumes allow us to share data between containers in the same pod

M.Romdhani, 2020

4

4

Volumes vs. Persistent Volumes

Managing stateful applications

■ Volumes and Persistent Volumes are related, but very different!

■ Volumes:

- appear in Pod specifications (we'll see that in a few slides)
- do not exist as API resources (cannot do `kubectl get volumes`)

■ Persistent Volumes:

- are API resources (can do `kubectl get persistentvolumes`)
- correspond to concrete volumes (e.g. on a SAN, EBS, etc.)
- cannot be associated with a Pod directly; but through a Persistent Volume Claim

M.Romdhani, 2020

5

5

Adding a volume to a Pod

Managing stateful applications

■ We add a volume to Pod manifest

- This will mount that volume in a container in the Pod

■ By default, this volume will be an `emptyDir` (an empty directory)

- It will "shadow" the directory where it's mounted

■ We can add the volume in two places:

- at the Pod level (to declare the volume)
- at the container level (to mount the volume)

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-with-volume
spec:
  volumes:
    - name: www
  containers:
    - name: nginx
      image: nginx
      volumeMounts:
        - name: www
          mountPath: /usr/share/nginx/html/
```

M.Romdhani, 2020

6

6

Sharing a volume between two containers

Managing stateful applications

- The volume `www` is added at the pod level
- We have 2 containers
 - Nginx mounts `www` to `/usr/share/nginx/html`
 - Git mounts `www` to `/www/`
- As a result, Nginx now serves this website

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-with-git
spec:
  volumes:
  - name: www
  containers:
  - name: nginx
    image: nginx
    volumeMounts:
    - name: www
      mountPath: /usr/share/nginx/html/
  - name: git
    image: alpine
    command: [ "sh", "-c", "apk add git && git clone https://github.com/octocat/Spoon-Knife /www" ]
    volumeMounts:
    - name: www
      mountPath: /www/
  restartPolicy: OnFailure
```

M.Romdhani, 2020

7

7

Init Containers

Managing stateful applications

- In the previous example, there is a short period of time during which the website is not available
 - because the git container hasn't done its job yet !
- We can define containers that should execute before the main ones
 - They will be executed in order (instead of in parallel)
 - They must all succeed before the main containers are started
- Other uses of init containers
 - Load content
 - Generate configuration (or certificates)
 - Database migrations
 - Waiting for other services to be up

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-with-init
spec:
  volumes:
  - name: www
  containers:
  - name: nginx
    image: nginx
    volumeMounts:
    - name: www
      mountPath: /usr/share/nginx/html/
  initContainers:
  - name: git
    image: alpine
    command: [ "sh", "-c", "apk add git && git clone https://github.com/octocat/Spoon-Knife /www" ]
    volumeMounts:
    - name: www
      mountPath: /www/
```

M.Romdhani, 2020

8

8

ConfigMaps, Secrets, and StatefulSets

9

Configuration Pattern

Managing stateful
applications

- **Kubernetes has an integrated pattern for decoupling configuration from application or container**
 - This pattern makes use of two Kubernetes components:
 - **ConfigMaps**
 - **Secrets**

M.Romdhani, 2020

10

10

ConfigMaps

Managing stateful applications

- Externalized data stored within kubernetes.
- Can be referenced through several different means:
 - environment variable
 - a command line argument (via env var)
 - injected as a file into a volume mount
- Can be created from a manifest, literals, directories, or files directly.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: manifest-example
data:
  state: Belgium
  city: Brussels
  content: |
    Look at this,
    its multiline!
```

- Imperative style:

- \$ kubectl create configmap literal-example --from-literal="city=Brussels" --from-literal=state=Belgium
- \$ kubectl create configmap file-example --from-file=cm/city --from-file=cm/state

M.Romdhani, 2020

11

11

Secrets

Managing stateful applications

- Functionally identical to a ConfigMap.
- Stored as **base64 encoded content**.
- Encrypted at rest within etcd (if configured!).
- Stored on each worker node in tmpfs directory.
- Ideal for username/passwords, certificates or other sensitive information that should not be stored in a container.

M.Romdhani, 2020

12

12

Secrets

Managing stateful applications

■ type: There are three different types of secrets within Kubernetes:

- **docker-registry** - credentials used to authenticate to a container registry
- **generic/Opaque** - literal values from different sources
- **tls** - a certificate based secret

```
apiVersion: v1
kind: Secret
metadata:
  name: manifest-secret
type: Opaque
data:
  username: S3ViZXJuZXRlcw==
  password: cGFzc3dvcmQ=
```

■ data: Contains key-value pairs of base64 encoded content.

■ Imperative style:

- `$ kubectl create secret generic literal-secret --from-literal=username=administrator --from-literal=password=password`
- `kubectl create secret generic file-secret --from-file=secret/username --from-file=secret/password`

M.Romdhani, 2020

13

13

Injecting ConfigMaps and Secrets

Managing stateful applications

■ Injecting as environment variable

ConfigMap

```
apiVersion: batch/v1
kind: Job
metadata:
  name: cm-env-example
spec:
  template:
    spec:
      containers:
        - name: mypod
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args: ["printenv CITY"]
          env:
            - name: CITY
              valueFrom:
                configMapKeyRef:
                  name: manifest-example
                  key: city
          restartPolicy: Never
```

Secret

```
apiVersion: batch/v1
kind: Job
metadata:
  name: secret-env-example
spec:
  template:
    spec:
      containers:
        - name: mypod
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args: ["printenv USERNAME"]
          env:
            - name: USERNAME
              valueFrom:
                secretKeyRef:
                  name: manifest-example
                  key: username
          restartPolicy: Never
```

M.Romdhani, 2020

14

14

Injecting ConfigMaps and Secrets

Managing stateful applications

■ Injecting in a command

```

apiVersion: batch/v1
kind: Job
metadata:
  name: cm-env-example
spec:
  template:
    spec:
      containers:
        - name: mypod
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args: ["echo Hello ${CITY}!"]
          env:
            - name: CITY
              valueFrom:
                configMapKeyRef:
                  name: manifest-example
                  key: city
          restartPolicy: Never

```

ConfigMap

```

apiVersion: batch/v1
kind: Job
metadata:
  name: secret-env-example
spec:
  template:
    spec:
      containers:
        - name: mypod
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args: ["echo Hello ${USERNAME}!"]
          env:
            - name: USERNAME
              valueFrom:
                secretKeyRef:
                  name: manifest-example
                  key: username
          restartPolicy: Never

```

Secret

M.Romdhani, 2020

15

15

Injecting ConfigMaps and Secrets

Managing stateful applications

■ Injecting as a Volume

```

apiVersion: batch/v1
kind: Job
metadata:
  name: cm-vol-example
spec:
  template:
    spec:
      containers:
        - name: mypod
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args: ["cat /myconfig/city"]
          volumeMounts:
            - name: config-volume
              mountPath: /myconfig
          restartPolicy: Never
      volumes:
        - name: config-volume
          configMap:
            name: manifest-example

```

ConfigMap

```

apiVersion: batch/v1
kind: Job
metadata:
  name: secret-vol-example
spec:
  template:
    spec:
      containers:
        - name: mypod
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args: ["cat /mysecret/username"]
          volumeMounts:
            - name: secret-volume
              mountPath: /mysecret
          restartPolicy: Never
      volumes:
        - name: secret-volume
          secret:
            secretName: manifest-example

```

Secret

M.Romdhani, 2020

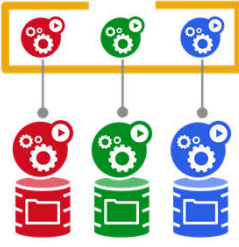
16

16

Statefulsets

Managing stateful applications

- Tailored to managing Pods that must persist or maintain state.
- Pod lifecycle will be ordered and follow consistent patterns.
 - Assigned a unique ordinal name following the convention of '<statefulset name>-<ordinal index>'.
- At a first glance, they look like deployments, but they have some significant differences
 - a stateful set defines a pod spec and a number of replicas R
 - it will make sure that R copies of the pod are running
 - updating the pod spec will cause a rolling update to happen



17

M.Romdhani, 2020

17

Statefulsets

Managing stateful applications

- StatefulSet Deployments provide:
 - Stable, unique network identifiers
 - Stable, persistent storage
 - Ordered, graceful deployment and scaling
 - Ordered, automated rolling updates

18

M.Romdhani, 2020

18

StatefulSet features

Managing stateful applications

- Pods in a stateful set are **numbered** (from 0 to R-1) and ordered
- They are **started and updated in order** (from 0 to R-1)
- A pod is started (or updated) only when the previous one is ready
- They are **stopped in reverse order** (from R-1 to 0)
- **Each pod know its identity** (i.e. which number it is in the set)
- Each pod can discover the IP address of the others easily
- The pods can persist data on attached volumes

M.Romdhani, 2020

19

19

Why not using just Volumes ?

Managing stateful applications

- **We can attach volumes to pods and deployments, but there are some shortcomings**
 - If a Deployment uses a volume, all replicas end up using the same volume
 - That volume must then support concurrent access
 - Their lifecycle (creation, deletion...) is managed outside of the Kubernetes API
- **What we really need is a way for each replica to have its own volume**
 - **The Pods of a Stateful set can have individual volumes**
 - in a Stateful set with 3 replicas, there will be 3 volumes
 - This introduces a bunch of new Kubernetes resource types: Persistent Volumes, Persistent Volume Claims, Storage Classes

M.Romdhani, 2020

20

20

StatefulSet manifest

Managing stateful applications

■ serviceName:

- The name of the associated headless service; or a service without a **ClusterIP**.
- serviceName is first unique aspect of statefulsets
 - maps to the name of a service you create along with the StatefulSet that helps provide pod network identity
 - Headless service, or a service without a ClusterIP. No NAT'ing. No load balancing.
 - It will contain a list of the pod endpoints and provides a mechanism for unique DNS names

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: sts-example
spec:
  replicas: 2
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: stateful
  serviceName: app
  updateStrategy:
    type: RollingUpdate
    rollingUpdate:
      partition: 0
  template:
    <pod template>
```

■ revisionHistoryLimit:

- The number of previous iterations of the StatefulSet to retain (as in Deployments and DaemonSets).

M.Romdhani, 2020

21

21

Statefulset, a recap

Managing stateful applications

- A Stateful sets manages a number of identical pods (like a Deployment)
- These pods are numbered, and started/upgraded/stopped in a specific order
- These pods are aware of their number (e.g., #0 can decide to be the primary, and #1 can be secondary)
- These pods can find the IP addresses of the other pods in the set (through a headless service)
- These pods can each have their own persistent storage (Deployments cannot do that)

M.Romdhani, 2020

22

22

Headless service

Managing stateful applications

- A headless service is a service with a service IP but **instead of load-balancing it will return the IPs of our associated Pods**.
- This allows us to interact directly with the Pods instead of a proxy. It's as simple as specifying None for `.spec.clusterIP` and can be utilized with or without selectors - you'll see an example with selectors in a moment.

```
apiVersion: v1
kind: Service
metadata:
  name: my-headless-service
spec:
  clusterIP: None # <--
  selector:
    app: test-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 3000
```

M.Romdhani, 2020

23

23

Headless service Example

Managing stateful applications

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: api-deployment
  labels:
    app: api
spec:
  replicas: 5
  selector:
    matchLabels:
      app: api
  template:
    metadata:
      labels:
        app: api
    spec:
      containers:
        - name: api
          image: eddiehale/hellonodeapi
          ports:
            - containerPort: 3000
```

Deployment
of 5 pods

```
apiVersion: v1
kind: Service
metadata:
  name: normal-service
spec:
  selector:
    app: api
  ports:
    - protocol: TCP
      port: 80
      targetPort: 3000
```

A Regular Service

```
apiVersion: v1
kind: Service
metadata:
  name: headless-service
spec:
  clusterIP: None # <-- Don't forget!!
  selector:
    app: api
  ports:
    - protocol: TCP
      port: 80
      targetPort: 3000
```

A Headless Service

- Hop in the cluster and run `nslookup`
 - `nslookup normal-service`, will show one entry
 - `nslookup headless-service` will show five entries

M.Romdhani, 2020

24

24

Storage classes

25

What are Storage Classes ?

Managing stateful
applications

- **Storage classes are an abstraction on top of a external storage resource that has dynamic provisioning capability**
 - Work hand-in-hand with the external storage system to enable dynamic provisioning of storage by eliminating the need for the cluster admin to pre-provision a volume
- **A StorageClass provides a way for administrators to describe the “classes” of storage they offer.**
 - Each StorageClass contains the fields provisioner, parameters, and reclaimPolicy, which are used when a PersistentVolume belonging to the class needs to be dynamically provisioned.

M.Romdhani, 2020

26

26

StorageClass manifest

Managing stateful applications

- **provisioner**: Defines the 'driver' to be used for provisioning of the external storage.
- **parameters**: A hash of the various configuration parameters for the provisioner.
- **reclaimPolicy**: The behaviour for the backing storage when the PVC is deleted.
 - **Retain** - manual clean-up
 - **Delete** - storage asset deleted by provider

```

apiVersion: v1
kind: PersistentVolume
metadata:
  name: task-pv-volume
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 1Mi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/C/tmp/data"

```

M.Romdhani, 2020

27

27

Available Storage Options

Managing stateful applications

- **Kubernetes storage options :**
 - Node Local storage : **emptyDir** (initially empty, the same lifecycle as the pod), **hostPath** (file or directory from the host node's filesystem into your pod)
 - File-sharing types such as nfs
 - Cloud provider-specific types like awsElasticBlockStore, azureDisk, VsphereVolume, PortworxVolume, gcePersistentDisk
 - Distributed file system types, for example glusterfs or cephfs
- **More are supported through plugins**

M.Romdhani, 2020

28

28

Persistent Volumes (PVs) and Persistent Volume Claims (PVCs)

29

What is a Persistent Volume

Managing stateful
applications

- A **PersistentVolume (PV)** represents a storage resource.
 - Containers are ephemeral constructs. Any changes to the running container is **lost** when the container stops running. PV are there to persist data for containers and Pods.
- PVs are a cluster wide resource linked to a backing storage provider: NFS, GCEPersistentDisk, RBD etc.
 - Generally provisioned by an administrator.
- Their lifecycle is handled independently from a pod
- CANNOT be attached to a Pod directly. Relies on a **PersistentVolumeClaim (PVC)**
 - The way a user consumes a PV is by creating a PVC.

M.Romdhani, 2020

30

30

PVs vs. PVCs

Managing stateful applications

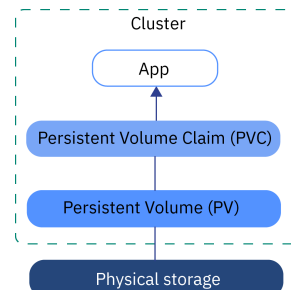
■ Persistent Volume (PV) –

- It is a piece of network storage that has been provisioned by the administrator. It's a resource in the cluster which is independent of any individual pod that uses the PV.

■ Persistent Volume Claim (PVC)

- It is a request for storage by a user that can be fulfilled by a PV.

■ PVs and PVCs are independent from Pod lifecycles and preserve data through restarting, rescheduling, and even deleting Pods.



M.Romdhani, 2020

31

31

Configure a Pod to Use a PersistentVolume for Storage

Managing stateful applications

■ Step 1 - Create a PersistentVolume

- The configuration file specifies that the volume is at /mnt/data on the cluster's Node
- View information about the PersistentVolume:

```
kubectl get pv task-pv-volume
```

```

apiVersion: v1
kind: PersistentVolume
metadata:
  name: task-pv-volume
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 1Mi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/C/tmp/data"
  
```

■ Step 2 - Create a PersistentVolumeClaim

- This **PersistentVolumeClaim** requests a volume of at 1 megabytes that can provide read-write access for at least one Node
- Look again at the PersistentVolume:

```
kubectl get pv task-pv-volume
```

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: task-pv-claim
spec:
  storageClassName: manual
  accessModes:
    - ReadWriteOnce
resources:
  requests:
    storage: 1Mi
  
```

M.Romdhani, 2020

32

Configure a Pod to Use a PersistentVolume for Storage

Managing stateful applications

Step 3 - Create a Pod

- Notice that the Pod's configuration file specifies a PersistentVolumeClaim, but it does not specify a PersistentVolume. From the Pod's point of view, the claim is a volume.

- Verify that the container in the Pod is running:

```
kubectl get pod task-pv-pod
```

- Initialize the index.html page from within C:\tmp\data

```
echo "Bonjour, Bonjour ..."
>/usr/share/nginx/html/index.html
```

- Check that /usr/share/nginx/html contains index.html having the right content.

- Get a shell to the container running in your Pod:

```
kubectl exec -it task-pv-pod -- /bin/bash
```

```
apiVersion: v1
kind: Pod
metadata:
  name: task-pv-pod
spec:
  volumes:
    - name: task-pv-storage
      persistentVolumeClaim:
        claimName: task-pv-claim
  containers:
    - name: task-pv-container
      image: nginx
      ports:
        - containerPort: 80
          name: "http-server"
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: task-pv-storage
```

M.Romdhani, 2020

33

33

Configure a Pod to Use a PersistentVolume for Storage

Managing stateful applications

Step 3 - Create a Pod (Continued)

- From within the pod shell, run the following commands


```
# Be sure to run these 3 commands inside the root shell that comes from
# running "kubectl exec" in the previous step
apt update
apt install -y curl
curl http://localhost/
```
- The curl output shows the text that you wrote to the **index.html** file on the hostPath volume (the string "Bonjour, Bonjour ...")

Step 4 - Clean Up

- Delete the Pod, the PersistentVolumeClaim and the PersistentVolume:


```
kubectl delete pod task-pv-pod
kubectl delete pvc task-pv-claim
kubectl delete pv task-pv-volume
```
- In the shell on your Node, remove the file and directory that you created:


```
rm /mnt/data/index.html
rmdir /mnt/data
```

M.Romdhani, 2020

34

34